


MADS - Deployment 1

It works on my machine

Raoul Grouls, 10 November 2025




Different types of data scientists

“I understand the algorithms, let someone else worry about implementation”

Different types of data scientists

- Writing python code that is acceptable for deployment
- Make your algorithm work on other machines beside your own laptop
- Collaborate with colleagues that build a frontend
- Use code that is published by others as a docker container



Docker for Data Scientists

Why containerization matters


Science is about reproducibility

Docker for Data Scientists: Getting Started

Why containerization matters for ML/AI projects

Promise of reproducibility across different machines

- Same environment
- Same dependencies
- Same behavior
- Isolation of dependencies from your personal workflow



Docker for Data Scientists: Getting Started


Why containerization matters for ML/AI projects

Easy sharing of complete project setups

- Less "it works on my machine"
- Fast setup on clean VMs
- Quick onboarding of new team members
(if they have enough RAM for Docker...)

Simplified deployment process


- From local to production with minimal changes



The "Works on My Machine" Paradox

Docker doesn't completely solve architecture differences:

- CPU Architecture Challenges
 - x86_64 (Intel/AMD) vs ARM64 (M1/M2 Macs)
 - Some packages aren't available for all architectures
 - Performance can vary significantly
- Common Issues
 - TensorFlow/PyTorch optimizations differ
 - C++ extensions might not compile
 - Binary dependencies might be architecture-specific
- Possible Solutions
 - Use --platform flag when building
 - Build multi-architecture images
 - Document architecture requirements clearly



Cross-Platform Docker: Using Platform Flags

Force AMD64 build

- docker build --platform linux/amd64 -t mymodel:latest .

Force ARM64 build

- docker build --platform linux/arm64 -t mymodel:latest .

Run for a specific platform

- docker run --platform linux/amd64 my-image

Docker Fundamentals: Key Concepts

Core Components:

- **Dockerfile** : The build instructions for creating an **Image**. Similar to a recipe with step-by-step instructions
- **Image** : The blueprint/template for **Containers**. Like a snapshot that can create identical containers
- **Container** : A running instance with your application and dependencies. Think of it as a lightweight VM that shares the host OS kernel
- **Registry** : A repository for storing and sharing images. Can be public (Docker Hub) or private

Docker Fundamentals: Key Concepts

Dockerfile -> builds -> **Image** -> runs -> **Container**

Understanding the Dockerfile

FROM

The command FROM:

Docker checks your local machine:

- Looks in local image cache
- If found, uses cached image
- If not found, needs to download

If not found locally:

- Connects to registry (default: Docker Hub)
- Downloads the image in layers
- Saves in local cache for future use

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Understanding the Dockerfile

FROM python:3.12-slim

There are multiple python images, you can find them on
https://hub.docker.com/_/python/tags

- -alpine
 - Ultra minimal (50MB)
 - Based on alpine linux
 - Security focused
 - Harder to debug
- -slim
 - Minimal image (130MB)
 - Only essentials
- Good for simple webapps like fastapi when size matters but alpine is too minimal
- Python:3.x
 - Full image (920MB)
 - Includes common system libraries
 - Good for complex dependencies
 - Larger deployment size
- -slim-bullseye / -slim-bookworm
 - Bullseye = Debian 11
 - Bookworn = Debian 12

Prebuilt baseimages

- With the script found here <https://github.com/raoulg/minimal-torch-docker>
- I publish minimal docker containers on hub.docker
- <https://hub.docker.com/repository/docker/raoulgrouls/torch-python-slim/general>
- You can find small prebuilt containers for arm64/amd64, python-slim 3.11/3.12, PyTorch 2.3-2.7 and with/without uv
- FROM raoulgrouls/torch-python-slim:py3.12-torch2.7.1-arm64-uv0.8.13

Understanding the Dockerfile

WORKDIR /app

- Sets the working directory
- All subsequent commands run here
- Good practice for organization

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Understanding the Dockerfile

COPY

- Copies files from host to container
- Best practice: Copy requirements first
- Helps with layer caching

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Understanding the Dockerfile

RUN

```
RUN pip install -r requirements.txt
```

- Executes commands during build
- Creates a new layer in the image
- Use && for multiple commands to reduce layers

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Understanding the Dockerfile

Layers

- Docker caches layers, such that they don't need to rebuild every time you run `build`
- Once docker detects that a layer has been changed, it will rebuild that layer, and all layers that come next
- In this example, once you change the requirements file, all subsequent layers will rebuild
- How would you improve this example?

```
1 FROM python:3.12-slim
2
3 WORKDIR /app
4
5 COPY requirements.txt requirements.txt
6 COPY src src
7
8 RUN pip install --no-cache-dir --upgrade -r requirements.txt
9
10 EXPOSE 8000
11
12 ENTRYPOINT ["uvicorn", "src.main:app", "--host", "0.0.0.0", "--port", "8000"]
```

CMD vs ENTRYPOINT

Making Your Container Executable

CMD:

- Can be easily overridden from command line
- If you provide runtime arguments, entire CMD is replaced
- Good for default behavior that might change

ENTRYPOINT:

- Runtime arguments are added to ENTRYPOINT arguments
- Harder to override (needs --entrypoint flag)
- Good for containers that always run the same command

CMD vs ENTRYPOINT

Making Your Container Executable

```
# Dockerfile
```

```
CMD ["python", "app.py"]
```

```
# Runtime - completely replaces CMD
```

```
docker run myimage echo "hello" # runs: echo "hello"
```

CMD vs ENTRYPOINT

Making Your Container Executable

```
# Dockerfile
```

```
ENTRYPOINT ["python"]
```

```
# Runtime - adds to ENTRYPOINT
```

```
docker run myimage app.py # runs: python app.py
```

```
docker run myimage -V # runs: python -V
```

CMD vs ENTRYPOINT

Making Your Container Executable

Best Practices:

- Use ENTRYPOINT for "main executable" that should always run
- Use CMD for arguments that might change
- Always use JSON array format ["command", "arg"]

```
ENTRYPOINT ["unicorn"]
```

```
CMD ["main:app", "--host", "0.0.0.0", "--port", "8000"]
```

Interactive Docker

Great for debugging

```
# Only overrides the CMD  
docker run -it my-image /bin/bash
```

```
# for Alpine-based images  
docker run -it my-image /bin/sh
```

```
# Ignores both ENTRYPOINT and CMD  
docker run -it --entrypoint /bin/bash my-image
```

Interactive Docker

Great for debugging

```
# Check installed Python packages  
pip list
```

```
# Test imports  
python  
>>> import numpy
```

```
# Check environment variables  
env
```

```
# Check file system  
ls -la  
cat requirements.txt
```

Makefile

Using Makefile to automate the docker commands

```
1 build:  
2     @echo "building docker image"  
3     docker build -t fastapi-app:latest .  
4  
5 run:  
6     docker run -p 8000:8000 fastapi-app:latest  
7     @echo "running on http://localhost:8000"  
8  
9 interactive:  
10    @echo "Entering container interactively"  
11    docker run -it --entrypoint /bin/bash fastapi-app:latest  
12
```